

Oil_Spill_Dataset_Analysis

July 23, 2024

```
[22]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
[2]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[3]: df = pd.read_csv('/content/drive/MyDrive/colab data file/oil_spill.csv')
df.head()
```

```
[3]:
```

	f_1	f_2	f_3	f_4	f_5	f_6	f_7	f_8	f_9	f_10	\
0	1	2558	1506.09	456.63	90	6395000	40.88	7.89	29780.0	0.19	
1	2	22325	79.11	841.03	180	55812500	51.11	1.21	61900.0	0.02	
2	3	115	1449.85	608.43	88	287500	40.42	7.34	3340.0	0.18	
3	4	1201	1562.53	295.65	66	3002500	42.40	7.97	18030.0	0.19	
4	5	312	950.27	440.86	37	780000	41.43	7.03	3350.0	0.17	

	...	f_41	f_42	f_43	f_44	f_45	f_46	f_47	f_48	\
0	...	2850.00	1000.00	763.16	135.46	3.73	0	33243.19	65.74	
1	...	5750.00	11500.00	9593.48	1648.80	0.60	0	51572.04	65.73	
2	...	1400.00	250.00	150.00	45.13	9.33	1	31692.84	65.81	
3	...	6041.52	761.58	453.21	144.97	13.33	1	37696.21	65.67	
4	...	1320.04	710.63	512.54	109.16	2.58	0	29038.17	65.66	

	f_49	target
0	7.95	1
1	6.26	0
2	7.84	1
3	8.07	1
4	7.35	0

[5 rows x 50 columns]

```
[4]: df.shape
```

```
[4]: (937, 50)
```

```
[5]: df.duplicated().sum()
```

```
[5]: 0
```

```
[6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 937 entries, 0 to 936
Data columns (total 50 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0    f_1     937 non-null    int64
 1    f_2     937 non-null    int64
 2    f_3     937 non-null    float64
 3    f_4     937 non-null    float64
 4    f_5     937 non-null    int64
 5    f_6     937 non-null    int64
 6    f_7     937 non-null    float64
 7    f_8     937 non-null    float64
 8    f_9     937 non-null    float64
 9    f_10    937 non-null    float64
10    f_11    937 non-null    float64
11    f_12    937 non-null    float64
12    f_13    937 non-null    float64
13    f_14    937 non-null    float64
14    f_15    937 non-null    float64
15    f_16    937 non-null    float64
16    f_17    937 non-null    float64
17    f_18    937 non-null    float64
18    f_19    937 non-null    float64
19    f_20    937 non-null    float64
20    f_21    937 non-null    float64
21    f_22    937 non-null    float64
22    f_23    937 non-null    int64
23    f_24    937 non-null    float64
24    f_25    937 non-null    float64
25    f_26    937 non-null    float64
26    f_27    937 non-null    float64
27    f_28    937 non-null    float64
28    f_29    937 non-null    float64
29    f_30    937 non-null    float64
30    f_31    937 non-null    float64
31    f_32    937 non-null    float64
```

```

32 f_33      937 non-null    float64
33 f_34      937 non-null    float64
34 f_35      937 non-null    int64
35 f_36      937 non-null    int64
36 f_37      937 non-null    float64
37 f_38      937 non-null    float64
38 f_39      937 non-null    int64
39 f_40      937 non-null    int64
40 f_41      937 non-null    float64
41 f_42      937 non-null    float64
42 f_43      937 non-null    float64
43 f_44      937 non-null    float64
44 f_45      937 non-null    float64
45 f_46      937 non-null    int64
46 f_47      937 non-null    float64
47 f_48      937 non-null    float64
48 f_49      937 non-null    float64
49 target    937 non-null    int64
dtypes: float64(39), int64(11)
memory usage: 366.1 KB

```

```

[7]: nv = df.isnull().sum()
      nv = nv[nv > 0]
      nv.sort_values(ascending=False)

```

```

[7]: Series([], dtype: int64)

```

```

[8]: df.duplicated().sum()

```

```

[8]: 0

```

```

[9]: df.describe()

```

```

[9]:
      count      f_1      f_2      f_3      f_4      f_5  \
count  937.000000  937.000000  937.000000  937.000000  937.000000
mean    81.588047  332.842049  698.707086  870.992209  84.121665
std     64.976730 1931.938570  599.965577  522.799325  45.361771
min       1.000000   10.000000   1.920000   1.000000   0.000000
25%      31.000000   20.000000  85.270000  444.200000  54.000000
50%      64.000000   65.000000 704.370000 761.280000 73.000000
75%     124.000000  132.000000 1223.480000 1260.370000 117.000000
max     352.000000 32389.000000 1893.080000 2724.570000 180.000000

      count      f_6      f_7      f_8      f_9      f_10  ...  \
count  9.370000e+02  937.000000  937.000000  937.000000  937.000000  ...
mean    7.696964e+05  43.242721   9.127887  3940.712914   0.221003  ...
std     3.831151e+06  12.718404   3.588878  8167.427625   0.090316  ...

```

min	7.031200e+04	21.240000	0.830000	667.000000	0.020000	...
25%	1.250000e+05	33.650000	6.750000	1371.000000	0.160000	...
50%	1.863000e+05	39.970000	8.200000	2090.000000	0.200000	...
75%	3.304680e+05	52.420000	10.760000	3435.000000	0.260000	...
max	7.131500e+07	82.640000	24.690000	160740.000000	0.740000	...

	f_41	f_42	f_43	f_44	f_45	\
count	937.000000	937.000000	937.000000	937.000000	937.000000	
mean	933.928677	427.565582	255.435902	106.112519	5.014002	
std	1001.681331	715.391648	534.306194	135.617708	5.029151	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	450.000000	180.000000	90.800000	50.120000	2.370000	
50%	685.420000	270.000000	161.650000	73.850000	3.850000	
75%	1053.420000	460.980000	265.510000	125.810000	6.320000	
max	11949.330000	11500.000000	9593.480000	1748.130000	76.630000	

	f_46	f_47	f_48	f_49	target
count	937.000000	937.000000	937.000000	937.000000	937.000000
mean	0.128068	7985.718004	61.694386	8.119723	0.043757
std	0.334344	6854.504915	10.412807	2.908895	0.204662
min	0.000000	2051.500000	35.950000	5.810000	0.000000
25%	0.000000	3760.570000	65.720000	6.340000	0.000000
50%	0.000000	5509.430000	65.930000	7.220000	0.000000
75%	0.000000	9521.930000	66.130000	7.840000	0.000000
max	1.000000	55128.460000	66.450000	15.440000	1.000000

[8 rows x 50 columns]

0.1 Selecting Categorical and numerical features

```
[10]: potential_categorical_features = []
threshold = 15

for column in df.columns:
    unique_values = df[column].nunique()
    if unique_values <= threshold:
        potential_categorical_features.append(column)
        print(f'Feature "{column}" has {unique_values} unique values and might_
        ↳ be categorical.')

if not potential_categorical_features:
    print("No potential categorical features found.")
else:
    print(f"Potential categorical features: {potential_categorical_features}")
```

Feature "f_22" has 9 unique values and might be categorical.
 Feature "f_23" has 1 unique values and might be categorical.

Feature "f_25" has 9 unique values and might be categorical.
 Feature "f_26" has 8 unique values and might be categorical.
 Feature "f_27" has 9 unique values and might be categorical.
 Feature "f_33" has 4 unique values and might be categorical.
 Feature "f_37" has 3 unique values and might be categorical.
 Feature "f_39" has 9 unique values and might be categorical.
 Feature "f_40" has 9 unique values and might be categorical.
 Feature "f_46" has 2 unique values and might be categorical.
 Feature "target" has 2 unique values and might be categorical.
 Potential categorical features: ['f_22', 'f_23', 'f_25', 'f_26', 'f_27', 'f_33', 'f_37', 'f_39', 'f_40', 'f_46', 'target']

```
[11]: cat_cols = ['f_22', 'f_23', 'f_25', 'f_26', 'f_27', 'f_33', 'f_37', 'f_39', 'f_40', 'f_46', 'target']

num_cols = [col for col in df.columns if col not in cat_cols]
```

```
[12]: df['f_40'].value_counts()
```

```
[12]: f_40
50      204
55      184
63      135
39      103
73       85
67       79
86       74
85       62
69       11
Name: count, dtype: int64
```

```
[13]: print(len(cat_cols))
print(len(num_cols))
```


```
11
39
```

0.2 Feature Selection

Feature selection is required here because we have 50 features which are unsuited for the model

```
[20]: corr = df.corr()

plt.figure(figsize=(30,30))
sns.heatmap(corr, annot=True, cmap= "Purples")
plt.show()
```



[21]: num_cols

[21]: ['f_1',
'f_2',
'f_3',
'f_4',
'f_5',
'f_6',
'f_7',
'f_8',

```
'f_9',  
'f_10',  
'f_11',  
'f_12',  
'f_13',  
'f_14',  
'f_15',  
'f_16',  
'f_17',  
'f_18',  
'f_19',  
'f_20',  
'f_21',  
'f_24',  
'f_28',  
'f_29',  
'f_30',  
'f_31',  
'f_32',  
'f_34',  
'f_35',  
'f_36',  
'f_38',  
'f_41',  
'f_42',  
'f_43',  
'f_44',  
'f_45',  
'f_47',  
'f_48',  
'f_49']
```

```
[23]: def correlation (dataset, threshold):  
    corr_col = set()  
    corr_matrix = dataset.corr()  
  
    for i in range(len(corr_matrix.columns)):  
        for j in range(i):  
            if abs(corr_matrix.iloc[i,j]) > threshold:  
                colname = corr_matrix.columns[i]  
                corr_col.add(colname)  
    return corr_col
```

```
[24]: High_corr_feat = correlation(df, 0.7)  
print(len(High_corr_feat))
```

```
[25]: High_corr_feat
```

```
[25]: {'f_11',  
      'f_12',  
      'f_13',  
      'f_14',  
      'f_15',  
      'f_16',  
      'f_18',  
      'f_20',  
      'f_21',  
      'f_24',  
      'f_25',  
      'f_29',  
      'f_30',  
      'f_32',  
      'f_34',  
      'f_35',  
      'f_36',  
      'f_38',  
      'f_39',  
      'f_40',  
      'f_42',  
      'f_43',  
      'f_44',  
      'f_48',  
      'f_49',  
      'f_6',  
      'f_9'}
```

Using Feature importance method of Random Forest for better decision

```
[26]: import pandas as pd  
      from sklearn.ensemble import RandomForestClassifier  
      from sklearn.model_selection import train_test_split
```

```
[27]: # Separate features and target  
x = df.drop('target', axis=1)  
y = df['target']  
  
# Split the dataset into training and testing sets  
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3,  
                                                    random_state=42)  
  
# Create and train the Random Forest model  
forest = RandomForestClassifier(random_state=42)  
forest.fit(x_train, y_train)
```



```
# Get feature importances
imp_feat = forest.feature_importances_
mean_importance = np.mean(imp_feat)
std_importance = np.std(imp_feat)
```

```
[28]: imp_feat
```

```
[28]: array([0.07332398, 0.01971647, 0.01506928, 0.04111053, 0.02176426,
            0.02163577, 0.02354695, 0.02084428, 0.01561991, 0.01614662,
            0.02156506, 0.01681376, 0.01984909, 0.01987316, 0.01594908,
            0.01767929, 0.02207773, 0.02288631, 0.01971611, 0.01317309,
            0.01514564, 0.00726373, 0.          , 0.01862542, 0.03859374,
            0.00502522, 0.00662783, 0.00911714, 0.01874776, 0.02655529,
            0.01861299, 0.01064635, 0.          , 0.0115403 , 0.01850997,
            0.02136448, 0.00108966, 0.01977912, 0.00359217, 0.00291329,
            0.03201321, 0.02162046, 0.02029819, 0.00998448, 0.0220041 ,
            0.00549814, 0.13096244, 0.02549836, 0.0200098 ])
```

```
[45]: imp_feat = forest.feature_importances_
indices = np.argsort(imp_feat)[::-1]
mean_importance = np.mean(imp_feat)
std_importance = np.std(imp_feat)


# Plot the feature importance of the forest

plt.figure(figsize=(9, 5))
plt.title("Feature Importance")

colormap = plt.cm.viridis
normalize = plt.Normalize(vmin=0, vmax=x_train.shape[1]-1)
colors = colormap(normalize(range(x_train.shape[1])))
plt.bar(range(x_train.shape[1]), imp_feat[indices], color = colors)

# Add mean feature importance line
plt.axhline(y=mean_importance, color='r', linestyle='--', label=f'Mean_
↳Importance: {mean_importance:.4f}')

plt.xticks(range(x_train.shape[1]), x_train.columns[indices], rotation=90)
plt.xlim([-1, x_train.shape[1]])
plt.ylim([0, 0.14])
plt.tight_layout()
plt.show()
```



```
[30]: features_below_mean = [x_train.columns[i] for i in range(len(imp_feat)) if
    ↪ imp_feat[i] < mean_importance]
print(len(features_below_mean))
print(features_below_mean)
```

32

```
['f_2', 'f_3', 'f_9', 'f_10', 'f_12', 'f_13', 'f_14', 'f_15', 'f_16', 'f_19',
'f_20', 'f_21', 'f_22', 'f_23', 'f_24', 'f_26', 'f_27', 'f_28', 'f_29', 'f_31',
'f_32', 'f_33', 'f_34', 'f_35', 'f_37', 'f_38', 'f_39', 'f_40', 'f_43', 'f_44',
'f_46', 'f_49']
```

```
[31]: features_below_mean = [x_train.columns[i] for i in range(len(imp_feat)) if
    ↪ imp_feat[i] < mean_importance]
features_below_mean
```

```
[31]: ['f_2',
'f_3',
'f_9',
'f_10',
'f_12',
'f_13',
'f_14',
'f_15',
'f_16',
'f_19',
'f_20',
'f_21',
```

```
'f_22',  
'f_23',  
'f_24',  
'f_26',  
'f_27',  
'f_28',  
'f_29',  
'f_31',  
'f_32',  
'f_33',  
'f_34',  
'f_35',  
'f_37',  
'f_38',  
'f_39',  
'f_40',  
'f_43',  
'f_44',  
'f_46',  
'f_49']
```

```
[32]: High_corr_feat
```

```
[32]: {'f_11',  
      'f_12',  
      'f_13',  
      'f_14',  
      'f_15',  
      'f_16',  
      'f_18',  
      'f_20',  
      'f_21',  
      'f_24',  
      'f_25',  
      'f_29',  
      'f_30',  
      'f_32',  
      'f_34',  
      'f_35',  
      'f_36',  
      'f_38',  
      'f_39',  
      'f_40',  
      'f_42',  
      'f_43',  
      'f_44',  
      'f_48',
```

```
'f_49',  
'f_6',  
'f_9']}
```

```
[46]: # Now selecting the common features from the High corr features and features_␣  
      ↪ below mean
```

```
unwanted_features = set(features_below_mean)  
High_corr_feat = set(High_corr_feat)  
  
drop_feat = unwanted_features & High_corr_feat # common elements selected  
drop_feat
```

```
[46]: {'f_12',  
      'f_13',  
      'f_14',  
      'f_15',  
      'f_16',  
      'f_20',  
      'f_21',  
      'f_24',  
      'f_29',  
      'f_32',  
      'f_34',  
      'f_35',  
      'f_38',  
      'f_39',  
      'f_40',  
      'f_43',  
      'f_44',  
      'f_49',  
      'f_9'}
```

```
[47]: orig_drop = {'f_12',  
                  'f_13',  
                  'f_14',  
                  'f_15',  
                  'f_16',  
                  'f_20',  
                  'f_21',  
                  'f_24',  
                  'f_29',  
                  'f_32',  
                  'f_34',  
                  'f_35',  
                  'f_38',
```

```
'f_40',  
'f_43',  
'f_44',  
'f_49',  
'f_9'}
```

```
[48]: print(len(orig_drp))
```

18

```
[49]: print(len(drop_feat))
```

19

```
[50]: diff = drop_feat - orig_drp  
print(diff)
```

{'f_39'}

```
[51]: # From the inference 'f_23' will also be dropped here
```

0.3 Treating Categorical columns

```
[52]: plt.figure(figsize=(20,20))  
for i in range(len(cat_cols)):  
    plt.subplot(4,3,i+1)  
    sns.countplot(y=df[cat_cols[i]])  
    plt.title(f'Countplot for {cat_cols[i]}')  
  
plt.show()
```



```
[ ]: ## Inference f_23 should be dropped
```

```
[ ]: df1 = df.copy()
      df2 = df.copy()
      df3 = df.copy()
```

```
[ ]: cat_cols.remove('f_23')
      df1.drop('f_23', axis=1, inplace=True)
      df1.drop(drop_feat, axis=1, inplace=True)
```

```
[ ]: df1.shape
```

```
[ ]: (937, 30)
```

```
[ ]: # Again selecting numcols and cat cols
```

```
[ ]: potential_categorical_features = []
threshold = 15

for column in df1.columns:
    unique_values = df1[column].nunique()
    if unique_values <= threshold:
        potential_categorical_features.append(column)
        print(f'Feature "{column}" has {unique_values} unique values and might
        ↳ be categorical.')

if not potential_categorical_features:
    print("No potential categorical features found.")
else:
    print(f"Potential categorical features: {potential_categorical_features}")
```

```
Feature "f_22" has 9 unique values and might be categorical.
Feature "f_25" has 9 unique values and might be categorical.
Feature "f_26" has 8 unique values and might be categorical.
Feature "f_27" has 9 unique values and might be categorical.
Feature "f_33" has 4 unique values and might be categorical.
Feature "f_37" has 3 unique values and might be categorical.
Feature "f_46" has 2 unique values and might be categorical.
Feature "target" has 2 unique values and might be categorical.
Potential categorical features: ['f_22', 'f_25', 'f_26', 'f_27', 'f_33', 'f_37',
'f_46', 'target']
```

```
[ ]: cat_cols = ['f_22', 'f_25', 'f_26', 'f_27', 'f_33', 'f_37', 'f_46', 'target']

num_cols = [col for col in df1.columns if col not in cat_cols]
print(len(cat_cols))
print(len(num_cols))
```

```
8
22
```

```
[ ]: print(len(num_cols))
print(num_cols)
```

```
22
['f_1', 'f_2', 'f_3', 'f_4', 'f_5', 'f_6', 'f_7', 'f_8', 'f_10', 'f_11', 'f_17',
'f_18', 'f_19', 'f_28', 'f_30', 'f_31', 'f_36', 'f_41', 'f_42', 'f_45', 'f_47',
'f_48']
```

```
[ ]:
```

0.4 Treating Numerical Columns

```
[ ]: print(len(cat_cols))
```

8

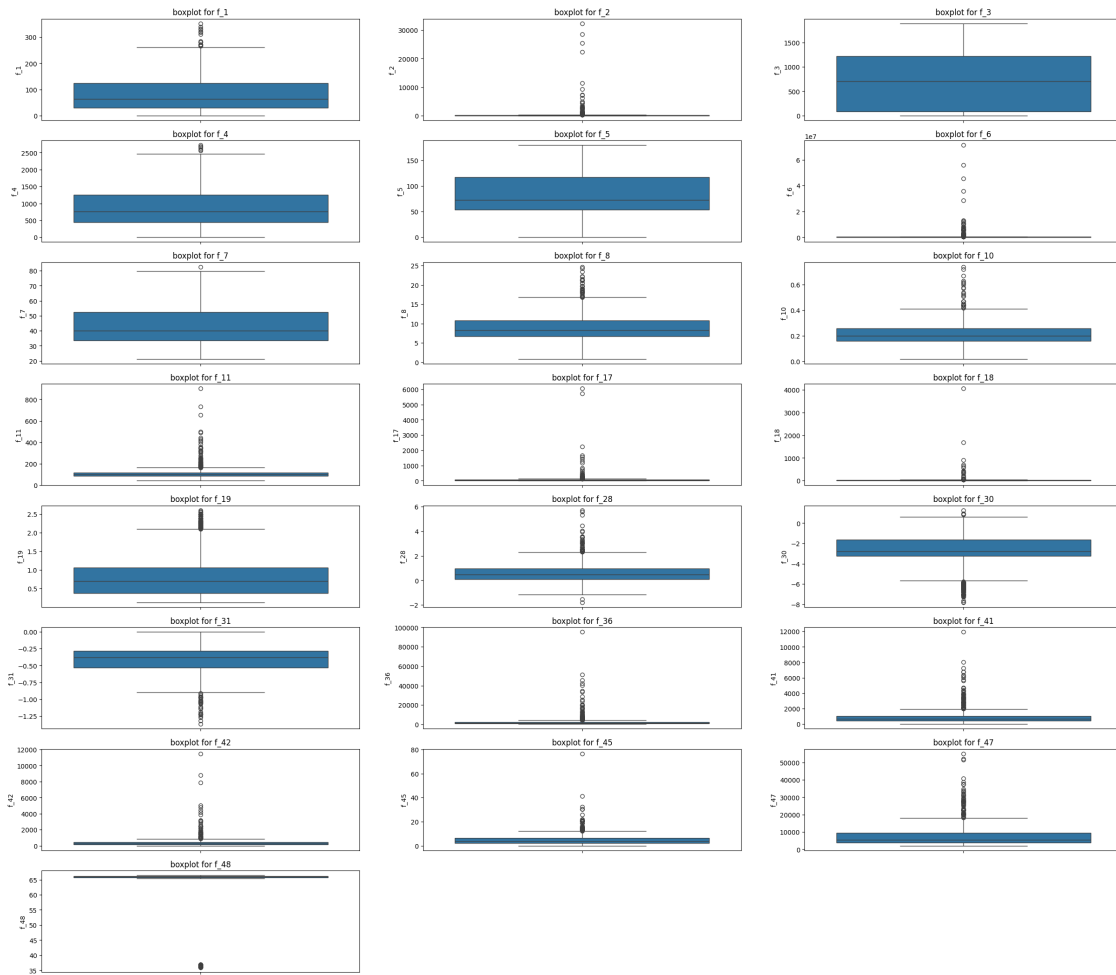
```
[ ]: num_cols
```

```
[ ]: ['f_1',  
      'f_2',  
      'f_3',  
      'f_4',  
      'f_5',  
      'f_6',  
      'f_7',  
      'f_8',  
      'f_10',  
      'f_11',  
      'f_17',  
      'f_18',  
      'f_19',  
      'f_28',  
      'f_30',  
      'f_31',  
      'f_36',  
      'f_41',  
      'f_42',  
      'f_45',  
      'f_47',  
      'f_48']
```

```
[ ]: print(len(num_cols))
```

22

```
[ ]: plt.figure(figsize=(30,30))  
      for i in range(len(num_cols)):  
          plt.subplot(9, 3, i+1)  
          sns.boxplot(y=df1[num_cols[i]])  
          plt.title(f'boxplot for {num_cols[i]}')  
  
      plt.show()
```

0.4.1 Outlier Treatment

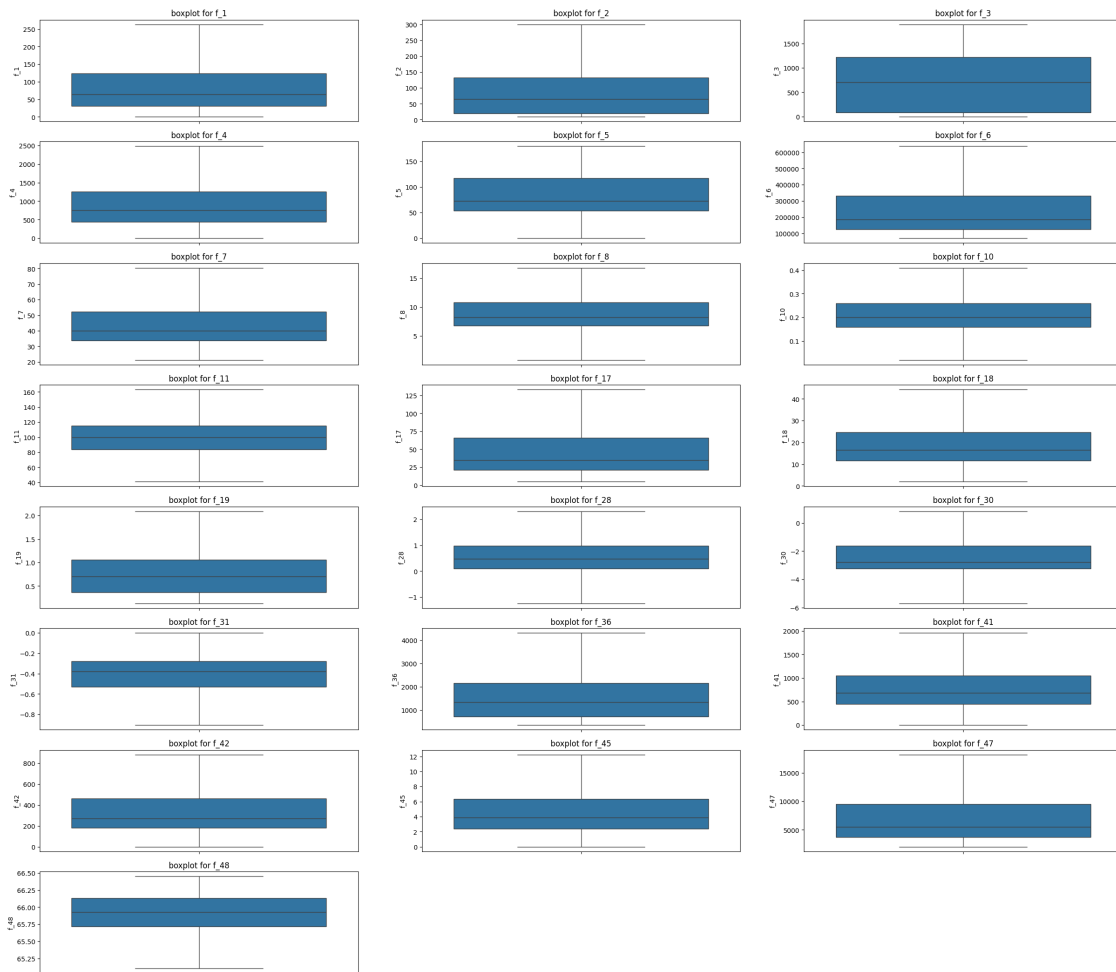
```
[ ]: num_cols
```

```
[ ]: ['f_1',
      'f_2',
      'f_3',
      'f_4',
      'f_5',
      'f_6',
      'f_7',
      'f_8',
      'f_10',
      'f_11',
      'f_17',
      'f_18',
      'f_19',
```

```
'f_28',  
'f_30',  
'f_31',  
'f_36',  
'f_41',  
'f_42',  
'f_45',  
'f_47',  
'f_48']
```

```
[ ]: for col in num_cols:  
    Q1 = df1[col].quantile(0.25)  
    Q3 = df1[col].quantile(0.75)  
    IQR = Q3 - Q1  
    lower_bound = Q1 - 1.5 * IQR  
    upper_bound = Q3 + 1.5 * IQR  
  
    df1[col] = df1[col].clip(lower=lower_bound, upper=upper_bound)
```

```
[ ]: plt.figure(figsize=(30,30))  
for i in range(len(num_cols)):  
    plt.subplot(9, 3, i+1)  
    sns.boxplot(y=df1[num_cols[i]])  
    plt.title(f'boxplot for {num_cols[i]}')  
  
plt.show()
```



```
[ ]: df1.shape
```

```
[ ]: (937, 30)
```

0.5 Splilt into x and y

```
[ ]: x = df1.drop('target', axis=1)
     y = df1['target']
```

```
[ ]: y.value_counts()
```

```
[ ]: target
0    896
1     41
Name: count, dtype: int64
```

0.6 Data Balancing

Data is Highly Imbalanced here so we need to balance it

```
[ ]: !pip install imblearn
```

Collecting imblearn

Downloading imblearn-0.0-py2.py3-none-any.whl (1.9 kB)

Requirement already satisfied: imbalanced-learn in

/usr/local/lib/python3.10/dist-packages (from imblearn) (0.10.1)

Requirement already satisfied: numpy>=1.17.3 in /usr/local/lib/python3.10/dist-packages (from imbalanced-learn->imblearn) (1.25.2)

Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.10/dist-packages (from imbalanced-learn->imblearn) (1.11.4)

Requirement already satisfied: scikit-learn>=1.0.2 in

/usr/local/lib/python3.10/dist-packages (from imbalanced-learn->imblearn) (1.2.2)

Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from imbalanced-learn->imblearn) (1.4.2)

Requirement already satisfied: threadpoolctl>=2.0.0 in

/usr/local/lib/python3.10/dist-packages (from imbalanced-learn->imblearn) (3.5.0)

Installing collected packages: imblearn

Successfully installed imblearn-0.0

```
[ ]: from imblearn.over_sampling import SMOTE
```

```
smote = SMOTE()
```

```
print(x.shape)
```

```
print(y.shape)
```

```
(937, 29)
```

```
(937,)
```

```
[ ]: x_res, y_res = smote.fit_resample(x, y)
```

```
[ ]: print(x_res.shape)
```

```
print(y_res.shape)
```

```
(1792, 29)
```

```
(1792,)
```

```
[ ]: y_res.value_counts()
```

```
[ ]: target
```

```
1    896
```

```
0    896
```

Name: count, dtype: int64

0.7 Split Data into train and test and standardize

```
[ ]: from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(x_res, y_res, test_size=0.
↳3, random_state=42)

print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```

(1254, 29)

(538, 29)

(1254,)

(538,)

```
[ ]: # Now Standardize the dataset
```

```
[ ]: from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

# Fit and transform the training data
x_train = scaler.fit_transform(x_train)

# Transform the test data using the same scaler
x_test = scaler.transform(x_test)
```

Creating evaluation metrics

```
[ ]: from sklearn.metrics import accuracy_score, precision_score, recall_score,
↳classification_report, confusion_matrix
```

```
[ ]: def eval_model(model, mname):
    model.fit(x_train, y_train)
    ypred = model.predict(x_test)

    accuracy = accuracy_score(y_test, ypred)
    precision = precision_score(y_test, ypred, average='binary')
    recall = recall_score(y_test, ypred, average='binary')
    class_report = classification_report(y_test, ypred)
    cm = confusion_matrix(y_test, ypred)
```

```

train_acc = model.score(x_train, y_train)
test_acc = model.score(x_test, y_test)

res_df = pd.DataFrame({'Train_Acc':train_acc,'Test_acc':test_acc,
↳ 'Accuracy': accuracy},
                        index=[mname])

# print('accuracy', accuracy)
# print('precision', precision)
# print('recall', recall)
print('classification_report', class_report)

return res_df

```

0.7.1 KNN m1

```

[ ]: from sklearn.neighbors import KNeighborsClassifier

knn_m1 = KNeighborsClassifier()
knn_m1.fit(x_train, y_train)

```

```

[ ]: KNeighborsClassifier()

```

```

[ ]: df_knn = eval_model(knn_m1, 'KNN')
df_knn

```

	classification_report		precision	recall	f1-score	support
	0	0.99	0.94	0.96		269
	1	0.94	0.99	0.96		269
	accuracy			0.96		538
	macro avg	0.96	0.96	0.96		538
	weighted avg	0.96	0.96	0.96		538

```

[ ]:      Train_Acc  Test_acc  Accuracy
KNN    0.980861  0.962825  0.962825

```

0.7.2 Logistic Regression m1

```

[ ]: from sklearn.linear_model import LogisticRegressionCV

log_m1 = LogisticRegressionCV()
log_m1.fit(x_train, y_train)

```

```

[ ]: LogisticRegressionCV()

```

```
[ ]: df_lg = eval_model(log_m1, 'Logistic Regression')
df_lg
```

classification_report		precision	recall	f1-score	support
	0	0.97	0.94	0.95	269
	1	0.94	0.97	0.96	269
accuracy			0.96		538
macro avg		0.96	0.96	0.96	538
weighted avg		0.96	0.96	0.96	538

```
[ ]:
Train_Acc Test_acc Accuracy
Logistic Regression 0.951356 0.95539 0.95539
```

```
[ ]:
```

0.7.3 Decision Tree m1

```
[ ]: from sklearn.tree import DecisionTreeClassifier

dt_m1 = DecisionTreeClassifier()
dt_m1.fit(x_train, y_train)
```

```
[ ]: DecisionTreeClassifier()
```

```
[ ]: df_dt = eval_model(dt_m1, 'Decision Tree') # Overfitting
df_dt
```

classification_report		precision	recall	f1-score	support
	0	0.96	0.96	0.96	269
	1	0.96	0.96	0.96	269
accuracy			0.96		538
macro avg		0.96	0.96	0.96	538
weighted avg		0.96	0.96	0.96	538

```
[ ]:
Train_Acc Test_acc Accuracy
Decision Tree 1.0 0.957249 0.957249
```

0.7.4 Random Forest m1

```
[ ]: from sklearn.ensemble import RandomForestClassifier

rf_m1 = RandomForestClassifier()
rf_m1.fit(x_train, y_train)
```

```
[ ]: RandomForestClassifier()
```

```
[ ]: df_rf = eval_model(rf_m1, 'Random Forest') # good results
df_rf
```

classification_report	precision	recall	f1-score	support
0	0.99	0.98	0.99	269
1	0.98	0.99	0.99	269
accuracy		0.99	538	
macro avg	0.99	0.99	0.99	538
weighted avg	0.99	0.99	0.99	538

```
[ ]: Train_Acc Test_acc Accuracy
Random Forest 1.0 0.986989 0.986989
```

0.7.5 Bagging Classifier m1

```
[ ]: from sklearn.ensemble import BaggingClassifier

bag_dt_m1 = BaggingClassifier(base_estimator=DecisionTreeClassifier(),
                               n_estimators=100, random_state=42)
bag_dt_m1.fit(x_train, y_train)
```

```
[ ]: BaggingClassifier(base_estimator=DecisionTreeClassifier(), n_estimators=100,
                      random_state=42)
```

```
[ ]: df_bag_dt = eval_model(bag_dt_m1, 'Bagging Decision Tree')
df_bag_dt
```

classification_report	precision	recall	f1-score	support
0	0.98	0.97	0.98	269
1	0.97	0.99	0.98	269
accuracy		0.98	538	
macro avg	0.98	0.98	0.98	538
weighted avg	0.98	0.98	0.98	538


```
[ ]:          Train_Acc  Test_acc  Accuracy
Bagging Decision Tree      1.0  0.975836  0.975836
```

```
[ ]: bag_knn_m1 = BaggingClassifier(base_estimator=KNeighborsClassifier(),
    ↪n_estimators=100, random_state=42)
bag_knn_m1.fit(x_train, y_train)
```

```
[ ]: BaggingClassifier(base_estimator=KNeighborsClassifier(), n_estimators=100,
    random_state=42)
```

```
[ ]: df_bag_knn = eval_model(bag_knn_m1, 'Bagging KNN')
df_bag_knn
```

classification_report			precision	recall	f1-score	support
	0	1.00	0.94	0.97		269
	1	0.94	1.00	0.97		269
accuracy				0.97		538
macro avg		0.97	0.97	0.97		538
weighted avg		0.97	0.97	0.97		538

```
[ ]:          Train_Acc  Test_acc  Accuracy
Bagging KNN      0.981659  0.966543  0.966543
```

```
[ ]: bag_rf_m1 = BaggingClassifier(base_estimator=RandomForestClassifier(),
    ↪n_estimators=100, random_state=42)
bag_rf_m1.fit(x_train, y_train)
```

```
[ ]: BaggingClassifier(base_estimator=RandomForestClassifier(), n_estimators=100,
    random_state=42)
```

```
[ ]: df_bag_rf = eval_model(bag_rf_m1, 'Bagging Random Forest') # overall good score
df_bag_rf
```

classification_report			precision	recall	f1-score	support
	0	0.99	0.98	0.98		269
	1	0.98	0.99	0.98		269
accuracy				0.98		538
macro avg		0.98	0.98	0.98		538
weighted avg		0.98	0.98	0.98		538

```
[ ]:
Train_Acc  Test_acc  Accuracy
Bagging Random Forest    0.999203  0.983271  0.983271
```

```
[ ]: bag_log_m1 = BaggingClassifier(base_estimator=LogisticRegressionCV(),
    ↪n_estimators=100, random_state=42)
bag_log_m1.fit(x_train, y_train)
```

```
[ ]: BaggingClassifier(base_estimator=LogisticRegressionCV(), n_estimators=100,
    random_state=42)
```

```
[ ]: df_bag_lg = eval_model(bag_log_m1, 'Bagging Logistic Regression')
df_bag_lg
```

classification_report		precision	recall	f1-score	support
0	0.97	0.93	0.95		269
1	0.94	0.97	0.95		269
accuracy			0.95		538
macro avg	0.95	0.95	0.95		538
weighted avg	0.95	0.95	0.95		538

```
[ ]:
Train_Acc  Test_acc  Accuracy
Bagging Logistic Regression    0.950558  0.951673  0.951673
```

```
[ ]: result_df = pd.concat([df_knn, df_lg, df_dt, df_rf, df_bag_dt, df_bag_knn,
    ↪df_bag_rf, df_bag_lg], axis=0).sort_values(by='Accuracy', ascending=False)
result_df
```

```
[ ]:
Train_Acc  Test_acc  Accuracy
Random Forest    1.000000  0.986989  0.986989
Bagging Random Forest    0.999203  0.983271  0.983271
Bagging Decision Tree    1.000000  0.975836  0.975836
Bagging KNN    0.981659  0.966543  0.966543
KNN    0.980861  0.962825  0.962825
Decision Tree    1.000000  0.957249  0.957249
Logistic Regression    0.951356  0.955390  0.955390
Bagging Logistic Regression    0.950558  0.951673  0.951673
```

So, As far as we see **Bagging Random Forest** has the higher results. we can apply hyperparameter tuning to it.

```
[ ]: from sklearn.model_selection import RandomizedSearchCV
from sklearn.ensemble import RandomForestClassifier
```

```
[ ]: # param_grid = {
#     'base_estimator__criterion': ['gini', 'entropy'],
#     'base_estimator__n_estimators': [50, 100, 200],
#     'base_estimator__max_features': ['auto', 'sqrt', 'log2'],
#     'base_estimator__max_depth': [None, 10, 20, 30],
#     'n_estimators': [10, 50, 100]
# }

# bag_rf = BaggingClassifier(base_estimator=RandomForestClassifier(),
#                             random_state=42)

[ ]: # rnd_ht = RandomizedSearchCV(bag_rf, param_grid, cv=5, scoring='accuracy')
# rnd_ht.fit(x_train, y_train)

[ ]: # dir(rnd_ht)

[ ]: # rnd_ht.best_params_

[ ]: # rnd_ht.best_score_
```

0.8 Model After Hyperparameter Tuning

```
[ ]: bag_rf_m1 = BaggingClassifier(base_estimator=RandomForestClassifier(),
#                                 random_state=42)
bag_rf_m1.fit(x_train, y_train)

[ ]: BaggingClassifier(base_estimator=RandomForestClassifier(), n_estimators=100,
                      random_state=42)

[ ]: RF = RandomForestClassifier(n_estimators=200, criterion='gini',
#                               max_features='log2', max_depth=20)

Bag_RF = BaggingClassifier(base_estimator=RF, n_estimators=50, random_state=42)
Bag_RF.fit(x_train, y_train)

[ ]: BaggingClassifier(base_estimator=RandomForestClassifier(max_depth=20,
#                                                         max_features='log2',
#                                                         n_estimators=200),
                      n_estimators=50, random_state=42)

[ ]: df_Bag_RF = eval_model(Bag_RF, 'Bagging Random Forest') # Select this model
df_Bag_RF
```

classification_report	precision	recall	f1-score	support
0	0.99	0.98	0.98	269

1	0.98	0.99	0.98	269
accuracy			0.98	538
macro avg	0.98	0.98	0.98	538
weighted avg	0.98	0.98	0.98	538

```
[ ]:          Train_Acc  Test_acc  Accuracy
Bagging Random Forest   0.999203  0.983271  0.983271
```

```
[ ]: Bag_RF.predict(x_test)
```

```
[ ]: array([0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1,
            0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0, 1,
            1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1,
            1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0,
            0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 0, 0, 0,
            1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 1,
            0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0,
            0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0,
            0, 1, 0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 1, 1, 1, 0, 1, 0, 0, 1, 0,
            0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1,
            0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1,
            0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 1,
            1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1,
            1, 1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0, 1,
            1, 1, 1, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
            0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1,
            0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 0, 0, 0, 1, 0, 1,
            0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1,
            0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 1,
            1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 0, 0, 0, 1, 1,
            0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0,
            1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1,
            0, 0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 1, 1, 1, 0, 1, 0, 0,
            1, 0, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1,
            1, 1, 1, 0, 0, 1, 0, 1, 1, 0])
```

0.9 Making new dataset with 20 samples

```
[ ]: new_data = df1.sample(20)
new_data.shape
```

```
[ ]: (20, 30)
```

```
[ ]: new_data
```

[]:	f_1	f_2	f_3	f_4	f_5	f_6	f_7	f_8	f_10	f_11	\
447	15.0	210	1111.80	1889.67	137	295312	31.80	8.43	0.27	88.40	
354	71.0	300	1166.82	1832.68	166	638670	40.51	11.27	0.28	116.30	
195	1.0	300	805.86	358.51	60	638670	24.65	6.74	0.27	163.25	
143	133.0	50	1687.60	1498.74	9	125000	37.34	7.30	0.20	63.10	
932	200.0	12	92.42	364.42	135	97200	59.42	10.34	0.17	110.00	
803	71.0	15	11.93	11.13	75	121500	44.40	10.02	0.23	105.30	
296	263.5	75	1117.59	842.47	55	187500	31.95	6.73	0.21	91.90	
601	33.0	12	122.83	578.33	76	97200	50.33	15.02	0.30	91.40	
781	49.0	27	11.74	554.74	68	218700	51.04	9.96	0.20	103.60	
635	67.0	12	356.25	614.92	95	97200	50.58	15.51	0.31	110.00	
734	2.0	98	386.31	292.22	110	638670	50.33	12.73	0.25	113.40	
239	110.0	170	1311.73	582.77	171	425000	30.55	6.98	0.23	95.30	
935	203.0	10	96.00	451.30	68	81000	59.90	15.01	0.25	97.50	
757	25.0	18	78.11	456.00	70	145800	48.94	7.79	0.16	91.80	
663	16.0	18	15.72	556.61	76	145800	69.22	10.98	0.16	100.90	
750	18.0	26	43.62	186.77	69	210600	51.35	6.47	0.13	86.50	
167	158.0	72	1706.17	1314.44	101	180000	34.01	9.15	0.27	71.70	
808	76.0	16	19.00	584.00	62	129600	50.12	7.80	0.16	112.30	
76	66.0	63	1563.17	1399.87	40	157500	37.70	8.34	0.22	97.20	
876	144.0	15	9.73	344.27	92	121500	54.80	16.47	0.30	103.40	

	...	f_33	f_36	f_37	f_41	f_42	f_45	f_46	f_47	f_48	\
447	...	0.0	2340	0.01	0.00	0.00	0.000	0	14192.31	65.105	
354	...	0.0	4320	0.01	1958.55	882.45	3.130	0	10219.05	66.340	
195	...	0.0	4320	0.00	1958.55	882.45	2.850	0	15749.24	65.610	
143	...	0.0	1620	0.01	509.90	304.14	2.460	0	3423.55	66.300	
932	...	0.0	540	0.01	381.84	254.56	4.500	0	2593.50	65.850	
803	...	0.0	450	0.00	569.21	180.00	5.420	0	13587.90	65.350	
296	...	0.0	1260	0.01	860.23	223.61	5.200	0	4866.52	65.870	
601	...	0.0	540	0.01	569.21	90.00	9.490	1	4513.34	66.140	
781	...	0.0	810	0.00	603.74	402.49	2.960	0	3772.43	66.070	
635	...	0.0	450	0.00	360.00	180.00	2.670	0	17382.04	66.280	
734	...	0.0	2160	0.00	1958.55	484.66	12.245	1	15413.42	65.870	
239	...	0.0	2790	0.01	912.41	608.28	2.210	0	7770.70	65.780	
935	...	0.0	540	0.01	402.49	180.00	4.470	0	2421.43	65.970	
757	...	0.0	720	0.00	853.81	180.00	12.200	0	2674.72	65.960	
663	...	0.0	630	0.00	569.21	201.25	4.610	0	3579.04	66.070	
750	...	0.0	990	0.00	1138.42	90.00	12.245	1	4390.92	65.590	
167	...	0.0	1710	0.01	873.21	206.16	6.250	0	4181.97	66.210	
808	...	0.0	720	0.01	649.00	127.28	10.200	1	3862.06	66.110	
76	...	0.0	900	0.01	430.12	320.16	2.070	0	5118.76	66.230	
876	...	0.0	630	0.01	450.00	270.00	3.120	0	2894.48	65.790	

	target
447	0
354	0

195	1
143	0
932	0
803	0
296	0
601	0
781	0
635	0
734	1
239	0
935	0
757	0
663	0
750	0
167	0
808	0
76	0
876	0

[20 rows x 30 columns]

0.9.1 Apply Preprocessing Steps

```
[ ]: new_data['target'].value_counts()
```

```
[ ]: target
0    18
1     2
Name: count, dtype: int64
```

```
[ ]: array1 = new_data['target'].to_numpy()
array1
```

```
[ ]: array([0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0])
```

```
[ ]: test_data = scaler.transform(new_data.drop('target', axis=1))
```

```
[ ]: predictions = Bag_RF.predict(test_data)

new_data['predictions'] = predictions
```

```
[ ]: new_data['predictions'].value_counts()
```

```
[ ]: predictions
0    18
1     2
Name: count, dtype: int64
```

```
[ ]: array2 = new_data['predictions'].to_numpy()  
array2
```

```
[ ]: array([0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0])
```

```
[ ]: array1 == array2
```

```
[ ]: array([ True,  True,  True,  True,  True,  True,  True,  True,  True,  
          True,  True,  True,  True,  True,  True,  True,  True,  True,  
          True,  True])
```

```
[ ]: # Calculate the number of correct predictions  
correct_predictions = np.sum(array1 == array2)  
  
# Calculate the accuracy percentage  
accuracy_percentage = (correct_predictions / len(array1)) * 100  
  
print(f'Accuracy: {accuracy_percentage:.2f}%')
```

Accuracy: 100.00%